

SUJITH S

AI Engineer | LLM & Multi-Agent Systems Developer
suijipk03@gmail.com • +91 6369817127 • [LinkedIn](#) • [GitHub](#) • [LeetCode](#)

EDUCATION

B.E – CS	Karpagam College of Engineering	2026	8.05 CGPA
HSC	KSV Matric Higher Secondary School	2022	8.4
SSLC	Mount Giris Matric Higher Secondary School	2020	9.4

TECHNICAL SKILLS

AI / ML	LLMs, Prompt Engineering, RAG (Retrieval-Augmented Generation), Multi-Agent Systems, Vector Embeddings, Semantic Search, NLP, Speaker Diarization
Frameworks	FastAPI, FastMCP, AgentSDK, OpenAI API (GPT-3.5 / GPT-4), NVIDIA NIM, React, Pandas, NumPy
Databases	MongoDB, Neo4j (Graph DB), ChromaDB (Vector Store), Redis
Cloud & DevOps	AWS Transcribe, Background Workers, Async Pipelines, GitHub, GitLab
Languages	Python, Java, C, JavaScript, HTML, SQL

EXPERIENCE

AI Engineer Intern Prayag.ai	May 2025 – Present
- Architected a FastMCP server exposing 15+ project management APIs as MCP tools, enabling an OpenAI-powered agent to dynamically select and invoke endpoints based on user intent — eliminating hardcoded routing logic entirely. - Designed a dual-source RAG pipeline combining ChromaDB vector embeddings (for documentation) and a Neo4j graph database (for codebase relationships), improving retrieval accuracy over flat-embedding approaches. - Built a 3-agent orchestration pipeline (intent classifier → hybrid retriever → response synthesizer) to deliver grounded, context-aware answers from heterogeneous knowledge sources. - Engineered an async document intelligence system using Redis-backed dual workers — one for AI embedding ingestion, one for Neo4j graph construction — processing 500+ legacy property documents with relationship extraction across configurable rule sets. - Integrated speaker diarization pipelines using Pyannote.audio and AWS Transcribe to extract per-speaker analytics from audio recordings for enterprise use cases. - Implemented OpenAI embeddings and vector storage workflows to power semantic search across large enterprise knowledge bases.	

KEY PROJECTS

PMT – MCP Conversational Agent

Aug 2025

Stack: Python, FastAPI, FastMCP, OpenAI API (GPT-4), AgentSDK

- **Designed and deployed a FastMCP server** that exposed 15+ REST API endpoints as structured MCP tools, enabling an LLM agent to dynamically resolve user intent to the correct API call without hardcoded routing.
- **Authored structured prompt schemas** defining tool parameters, usage intent, and response contracts — reducing agent hallucination and improving tool-selection accuracy across complex multi-step workflows.
- **Delivered a production chatbot** on top of a 40–60 API project management platform, making the entire application navigable through natural language.

Knowledge Base Chatbot

Oct 2025

Stack: Python, OpenAI API, ChromaDB, Neo4j, FastAPI

- **Architected a hybrid retrieval system** with two distinct data backends: ChromaDB vector store for unstructured documentation and Neo4j graph database for codebase — preserving structural and dependency relationships that flat embeddings cannot capture.
- **Built a 3-agent pipeline** (decider → data retriever → response generator) over 2 knowledge sources, enabling precise, citation-grounded answers to developer queries about internal systems.
- **Reduced retrieval ambiguity** by routing queries to the semantically appropriate source based on intent classification before retrieval.

Property Document Analyzer

Dec 2025

Stack: Python, FastAPI, Redis, Neo4j, OpenAI Embeddings

- **Built an async AI pipeline** processing 500+ legacy property documents using Redis-backed background workers, enabling non-blocking ingestion and graph construction at scale.
- **Separated concerns across dual workers** — one for document ingestion with OpenAI embeddings, one for Neo4j relationship graph construction — enabling parallel execution and faster end-to-end job completion.
- **Implemented configurable relationship-extraction rules** to identify legal and ownership links between documents, surfacing relationships invisible to keyword search.

EasyHire.ai – AI Hiring Assistant

Feb 2026

Stack: Python, FastAPI, NVIDIA NIM APIs (Embeddings + LLMs)

- **Automated resume screening and candidate shortlisting** using semantic similarity scoring with NVIDIA NIM embeddings, reducing manual review effort in high-volume hiring pipelines.
- **Deployed on scalable real-time infrastructure** via FastAPI, enabling concurrent processing of multiple candidate profiles with intelligent ranking output.

ACHIEVEMENTS & CERTIFICATIONS

- **LeetCode:** Solved 240+ problems; earned Pandas badge; 300+ problems solved across competitive platforms.
- **UI Design Competition — 2nd Place** (Kongu Engineering College, Mar 2023): Outperformed 50+ participants with user-centered interface designs.
- **Hackathon Team Lead** (Nov 2023): Led a 5-member team to build a web-based environmental sustainability solution; managed timelines and final presentation.